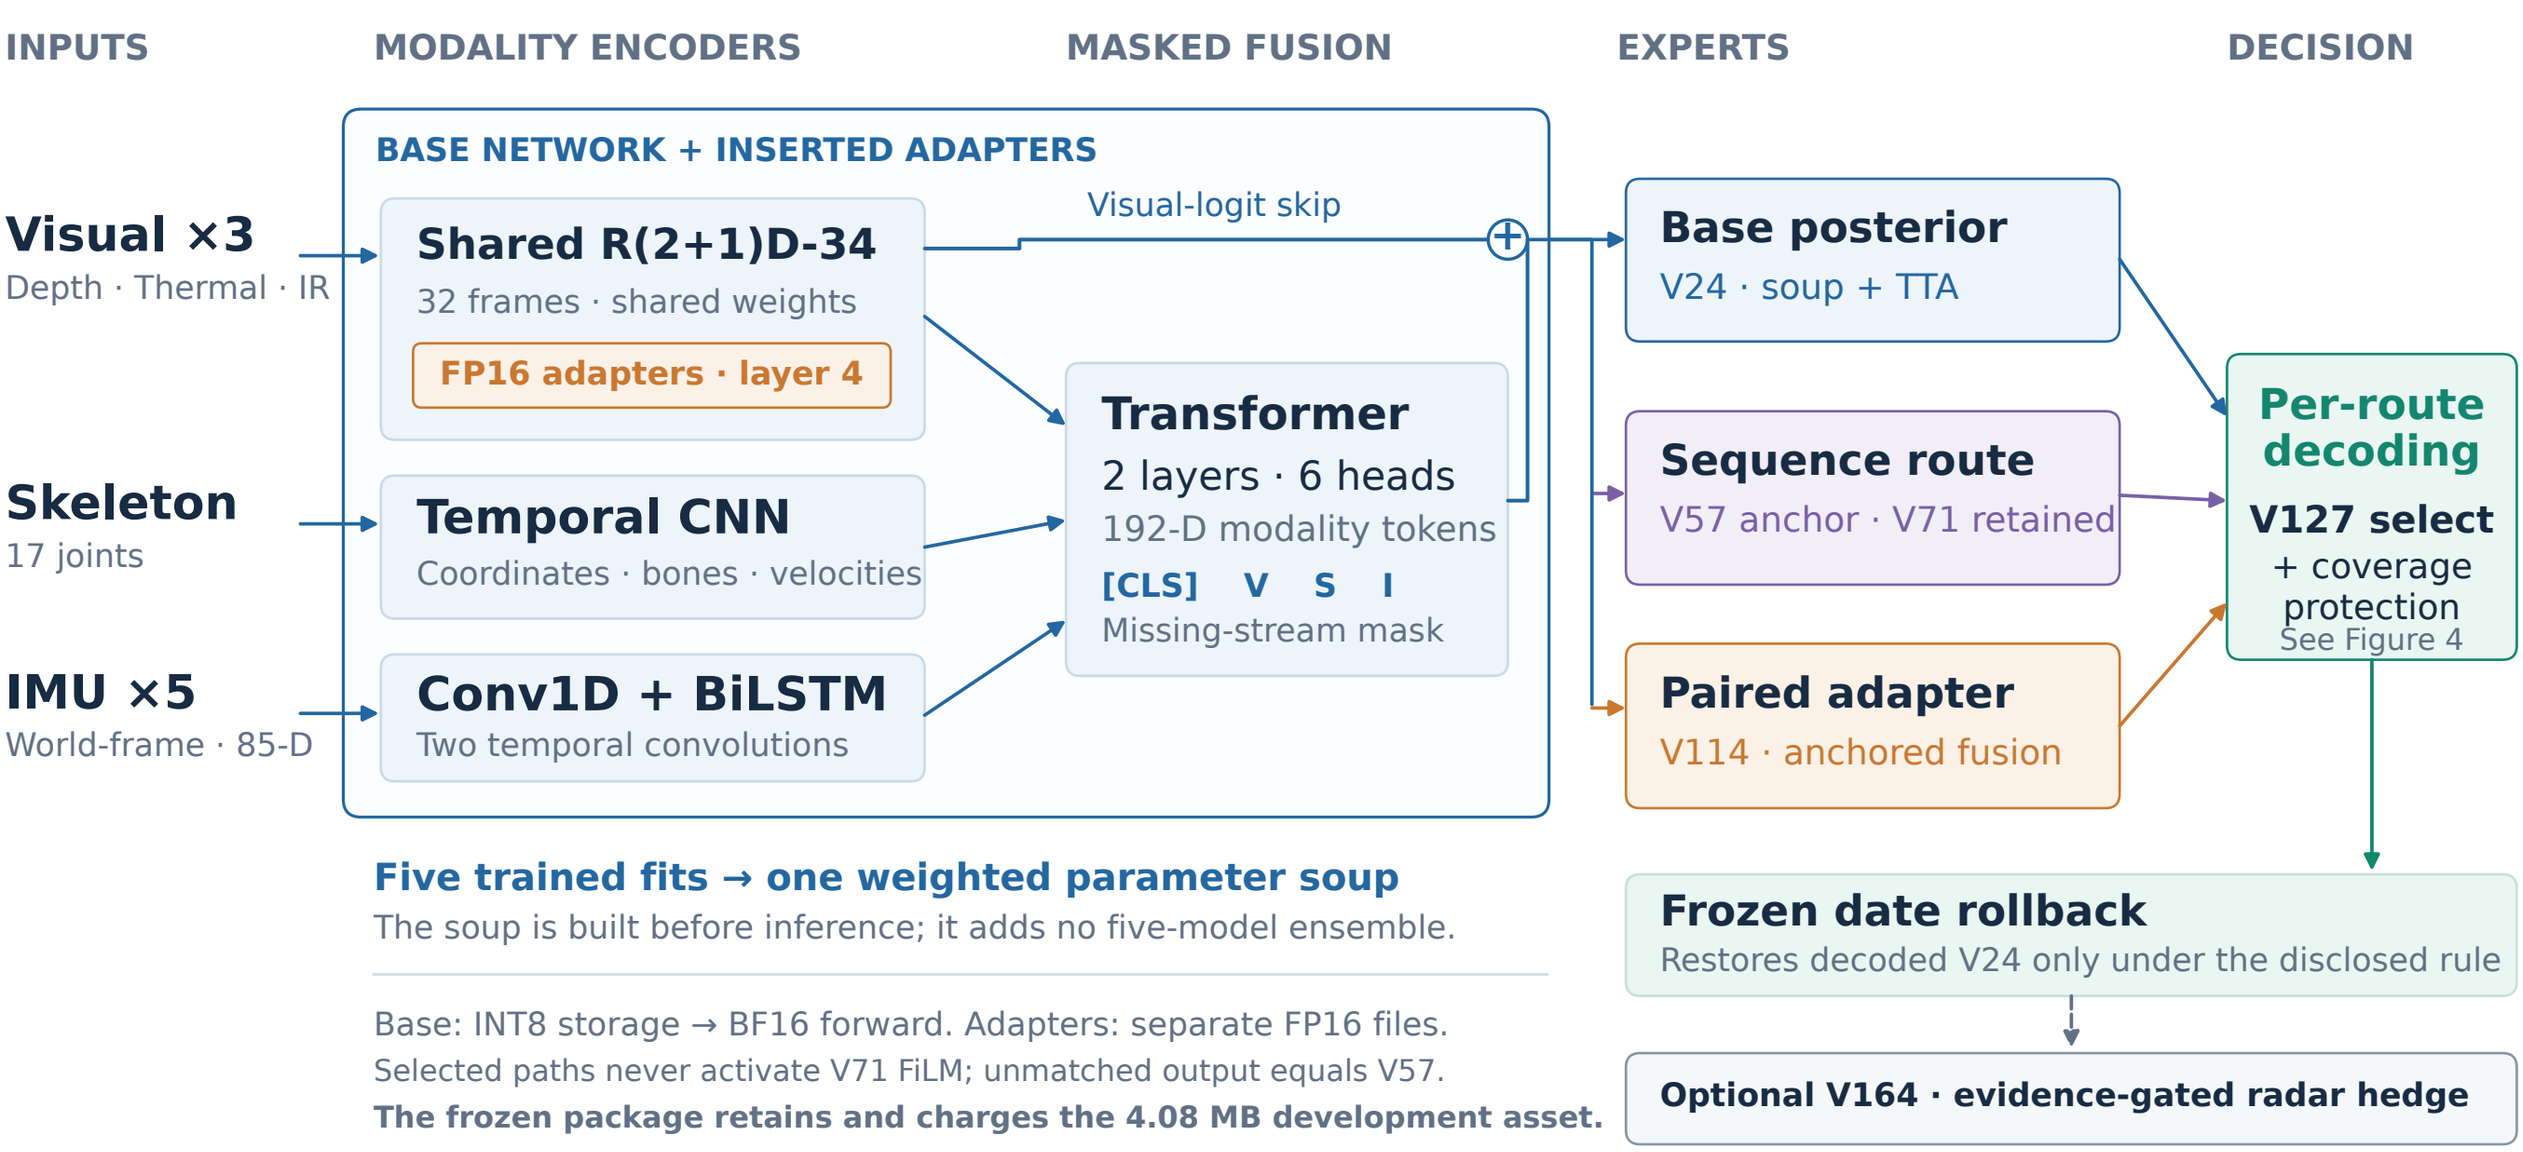
INPUTS
MODALITY ENCODERS
MASKED FUSION
EXPERTS
DECISION
BASE NETWORK + INSERTED ADAPTERS
Visual ×3
Depth · Thermal · IR
Skeleton
17 joints
IMU ×5
World-frame · 85-D
Shared R(2+1)D-34
32 frames · shared weights
FP16 adapters · layer 4
Temporal CNN
Coordinates · bones · velocities
Conv1D + BiLSTM
Two temporal convolutions
Visual-logit skip
Transformer
2 layers · 6 heads
192-D modality tokens
[CLS] V S I
Missing-stream mask
Base posterior
V24 · soup + TTA
Sequence route
V57 anchor · V71 retained
Paired adapter
V114 · anchored fusion
Per-route decoding
V127 select
+ coverage protection
See Figure 4
Frozen date rollback
Restores decoded V24 only under the disclosed rule
Optional V164 · evidence-gated radar hedge
Five trained fits → one weighted parameter soup
The soup is built before inference; it adds no five-model ensemble.
Base: INT8 storage → BF16 forward. Adapters: separate FP16 files.
Selected paths never activate V71 FiLM; unmatched output equals V57.
The frozen package retains and charges the 4.08 MB development asset.